

## RESEARCH ARTICLE

# When to Impute? Imputation before and during cross-validation

Byron C. Jaeger\*<sup>1</sup> | Nicholas J. Tierney<sup>2</sup> | Noah R. Simon<sup>3</sup>

<sup>1</sup>Department of Biostatistics, University of Alabama at Birmingham, Birmingham, Alabama

<sup>2</sup>Department of Econometrics and Business Statistics, Monash University, Melbourne, Victoria, Australia

<sup>3</sup>Department of Biostatistics, University of Washington, Seattle, Washington

**Correspondence**

\*Byron C. Jaeger. Email: bcjaeger@uab.edu

**Present Address**

327M Ryals Public Health Building 1665  
University Blvd Birmingham, Alabama  
35294-0022

Cross-validation (CV) is a common technique used to estimate generalization error for prediction models. For pipeline modeling algorithms (*i.e.*, modeling procedures with multiple steps), it has been recommended the *entire* sequence of steps be carried out during each replicate of CV to mimic the application of the entire pipeline to an external testing set. While theoretically sound, following this recommendation can lead to high computational costs when a pipeline modeling algorithm includes computationally expensive operations, *e.g.*, imputation of missing values. There is one exception: unsupervised variable selection (*i.e.*, ignoring the outcome) can be applied before conducting CV without incurring bias. However, it is unclear whether an unsupervised operation that modifies the training data (*i.e.*, imputation) rather than removing columns from the training data (*i.e.*, variable selection) can be applied prior to CV without causing model error estimates to become optimistically biased. We empirically assessed whether conducting unsupervised imputation prior to CV would result in biased estimates of generalization error or result in poorly selected tuning parameters and thus degrade the external performance of downstream models. Results show that despite optimistic bias, the reduced variance of imputation before CV compared to imputation during each replicate of CV leads to a lower overall root mean squared error for estimation of the true external  $R^2$  and the performance of models tuned using CV with imputation before versus during each replication is minimally different. In conclusion, unsupervised imputation before CV appears valid in certain settings and may be a helpful strategy that enables analysts to use more flexible imputation techniques without incurring high computational costs.

**KEYWORDS:**

Cross-validation; Missing data; Imputation; Machine learning;

## 1 | INTRODUCTION

In evaluating the performance of predictive modeling algorithm, it is understood that so-called training error (the predictive error measured on observations used to fit the model) is a poor proxy for generalization error (the performance of the model on

future, as-yet-unseen, observations).<sup>1</sup> The training error of a model will often be overly optimistic for the generalization error. As such, it is standard practice to use sample-splitting methods to estimate generalization error. These methods train and test models using separate datasets.  $v$ -fold Cross-validation (CV) is a common sample-splitting method that partitions a dataset into  $v$  non-overlapping subsets (*i.e.*, folds).<sup>2</sup> Each fold is then used as an internal assessment set for a modeling algorithm developed using data from the  $v - 1$  remaining folds. Aggregating errors from all  $v$  replications of this procedure provides an estimate of the modeling algorithm's generalization error, making  $v$ -fold CV an effective technique to 'tune' modeling algorithms (*i.e.*, select optimal values for parameters that govern the algorithm's fitting procedure).

Machine learning analyses often involve 'pipeline' modeling algorithms: multi-step modeling procedures that may include data pre-processing, predictor variable selection and/or transformation, model fitting, and ensembling.<sup>3</sup> For example, a pipeline may begin by centering and scaling predictor values, then filter out redundant correlated predictors, and finally fit a regression model to the remaining data. To estimate the generalization error of pipeline modeling algorithms using CV, it is recommended that the *entire sequence* of steps be carried out during each replicate of CV to mimic the application of the entire pipeline to an external testing set. However, it has been stated that unsupervised variable selection steps (*i.e.*, steps that ignore the outcome variable) can be applied before conducting CV without incurring bias.<sup>4</sup> Since unsupervised predictor variable selection does not involve outcome variables, it does not give the selected predictors an unfair advantage during CV.

Missing data (MD) occur frequently in machine learning analyses, and several learning algorithms (e.g., regression) are incompatible with MD. Imputation is a technique that replaces MD with estimated values, and is often among the most computationally expensive operations in pipeline modeling algorithms. For example, the `missForest` imputation algorithm may fit one random forest model for each column that contains MD.<sup>5</sup> Computational expense of applying `missForest` or other complex imputation strategies during each replicate of CV scales poorly and may lead analysts to prefer more convenient but less effective strategies to handle MD. A more computationally efficient approach would be to implement 'unsupervised imputation' (*i.e.*, imputing MD without accessing outcome information) *before* conducting CV. Ordering operations this way could result in substantially faster training and tuning costs for pipeline modeling algorithms because imputation is only performed once rather than once per CV fold. However, it is unclear whether an unsupervised operation that modifies the training data (*i.e.*, imputation) rather than removing columns from the training data (*i.e.*, variable selection) can be applied prior to CV without causing model error estimates to become overly optimistic.

The aim of this paper is to assess whether unsupervised imputation before versus during CV causes bias in estimation of a modeling pipeline's generalization error. We also investigate whether unsupervised imputation before versus during CV can result in poorly selected tuning parameters and thus degrade the external performance of downstream models. To achieve these aims, we conduct empirical studies of simulated and real data assessing whether unbiased generalization error estimates are obtained if unsupervised imputation is implemented before CV, a strategy we will refer to as  $I \rightarrow CV$ . We compare estimated

pipeline error according to  $I \rightarrow CV$  with estimated pipeline error when unsupervised imputation is applied *during each replicate* of CV, a strategy we will refer to as  $CV \cup I$ . All scripts involved in the current analysis are publicly available and all results are reproducible (See first author's GitHub). Our analysis also introduces and applies the `ipa` R package (imputation for predictive analytics), which provides functions to create single or multiple imputed training and testing sets for prediction modeling.

The rest of this manuscript is organized as follows. In Section 2, we discuss MD mechanisms and prevailing MD strategies for statistical inference and machine learning. In Section 3, we explicitly map the order of operations for  $CV \cup I$  and  $I \rightarrow CV$ . In section 4, we conduct a simulation study to assess empirical differences between  $CV \cup I$  and  $I \rightarrow CV$ . The two procedures are compared using real data in Section 5. Last, in Section 6, we organize the data from preceding sections to form recommendations for practitioners.

## 2 | MISSING DATA

### Missing data mechanisms

MD mechanisms were first formalized by Rubin,<sup>6</sup> who developed a framework to analyze MD that supposes each data point has some probability of being missing. If the probability of missingness is unrelated to the data (*i.e.*, all data are equally likely to be missing), then the data are missing completely at random (MCAR). When the probability of missingness is related to observed variables in the data (*i.e.*, all data within observed groups are equally likely to be missing), the data are missing at random (MAR). If the probability of missingness is determined by reasons that are unknown or unobserved, the data are missing not at random (MNAR). To illustrate, if a patient's data are missing because of clerical data entry error, the patient's data are MCAR. If instead a doctor chose not to measure the patient's labs because the patient was too young, the patient's data are MAR. If the patient missed the appointment because the patient was too sick, the patient's data are MNAR. In the context of statistical learning, previous findings have shown that when data are MNAR, imputation alone is often less effective than incorporating features that characterize missing patterns (*e.g.*, missingness incorporated as an attribute).<sup>7,8,9</sup> Since the primary aim of the current study is to assess the differences between two implementations of imputation (*i.e.*,  $I \rightarrow CV$  and  $CV \cup I$ ), we focus analyses on cases where data are MAR or MCAR.

### Missing data strategies for statistical inference

The primary objective for statistical inference in the presence of MD is to obtain valid test statistics for statistical hypotheses. Imputation to the mean and, more broadly, MD strategies that create a single imputed value, have been shown to increase type I errors (*i.e.*, rejecting a true null hypothesis) for inferential statistics by artificially reducing the variance of observed data and ignoring the uncertainty attributed to MD, respectively. Multiple imputation, a widely recommended strategy to handle MD for statistical inference, is capable of producing valid test statistics when data are MCAR or MAR because it can simultaneously

address these two shortcomings. The ‘accuracy’ of imputed values is not critical for the success of multiple imputation, given sufficient estimates of conditional distributions<sup>10</sup>. Instead, the consistency of the estimated covariance matrix for regression coefficients makes this strategy ideal for statistical inference.

### Missing data strategies for statistical prediction

The primary objective for statistical prediction in the presence of MD is to develop a prediction function that accurately generalizes to external data, which may or may not contain missing values (see Section 3.1). In contrast to statistical inference, single imputation is often used for prediction models.<sup>11</sup> Moreover, imputation strategies with greater accuracy often lead to better performance of downstream models (*i.e.*, models fitted to the imputed data). For example, Jerez et al. found single imputation using machine learning models provided superior downstream model prognostic accuracy compared to multiple imputation based on regression and expectation maximization.<sup>12</sup> The results of this analysis exemplify a perspective that will be taken throughout the current study. Namely, the authors treated imputation strategies as components of the modeling pipeline with parameters that can be ‘tuned’ in the same manner as a prediction model.

## 3 | ORDER OF OPERATIONS

In the context of statistical prediction, analysts usually work with a training set and an external testing set. A pipeline modeling algorithm developed with data from the training set can be externally validated using data from the testing set. Workflows to develop and validate a pipeline model may include three steps: (1) selection of pipeline parameter values (*i.e.*, parameters relevant to any operation in the pipeline, including data pre-processing), (2) developing a final model by training the modeling pipeline using the training data, and (3) externally validating the final model by assessing the accuracy of its predictions with the testing data. This workflow is described in **Figure 1**.

Pipeline parameter values may be set apriori or determined empirically (*i.e.*, tuned) using resampling, *e.g.*, by leveraging  $v$ -fold CV. We refer to the  $v - 1$  folds and 1 remaining fold used to internally train and test a modeling algorithm as **analysis** and **assessment** sets, respectively, to avoid notation abuse of the terms “training” and “testing.”<sup>13</sup> If CV is applied to facilitate selection of pipeline parameter values, it is critical that analysis data are separated from assessment data before any ‘learning’ is done. The entire *supervised* pipeline must be run using only the assessment data. This applies both to supervised data pre-processing steps (*e.g.*, selecting all variables with high correlation to the outcome) as well as supervised modeling procedures (*e.g.*, regression). ‘Data leakage’ can occur when outcome information from the assessment set is leveraged to modify the analysis set, *e.g.*, supervised variable selection is performed on a stacked set comprising analysis and assessment data, rather than just the assessment data.<sup>4</sup> There are a number of examples showing wildly optimistic estimates of generalization error because of

data leakage.<sup>14</sup> In scenarios with a larger number of features, even simple methodologies such as selecting those features with high individual correlation to the outcome can induce substantial overoptimism.

To remove any possibility of data leakage, all steps of the pipeline may be performed in analysis and assessment sets, separately, within each replicate of CV. For example, consider centering and scaling predictor variables such that they have zero mean and unit variance. As these operations do not involve the outcome, they are entirely unsupervised. Nevertheless, centering and scaling operations are usually completed in analysis and assessment sets, separately, during each replicate of CV. Specifically, the means and standard deviations are computed using the analysis data and then those values are applied to center and scale predictors in both the analysis and assessment sets. We refer to this traditional implementation of CV as  $CV \cup I$  and refer to the experimental implementation of CV (*i.e.*, one where unsupervised imputation occurs before CV begins) as  $I \rightarrow CV$  (**Figure 2**). Regardless of which implementation is applied, the output of CV is a set of pipeline parameter values and an estimate of generalization error. The pipeline parameter values are subsequently used to develop and validate a final prediction model using the full training set and testing set, respectively. Differences in parameter values selected by competing CV strategies (*i.e.*,  $CV \cup I$  or  $I \rightarrow CV$ ) may have measurable impact on the generalization error of their downstream models.

### 3.1 | Testing data

Ideally, external testing data will not contain MD, and imputation will not be necessary. However, If MD are present in the external testing data, additional steps may be taken to engage with them. One may impute missing values in the testing data using (1) only the training data, (2) only the testing data, or (3) using both training and testing data. It is common to use only the training data to impute missing values in the testing data. However, some imputation procedures can't be discretely separated into two steps, one that develops an imputation model and another that applies it to new data. For example, matrix decomposition methods such as `softImpute` merely take a matrix with missing entries and fill them in.<sup>15</sup> To apply these types of imputation procedures, approach (2) or (3) may be taken. Notably,  $CV \cup I$  uses only the analysis data to impute missing values in the assessment data (*i.e.*, approach 1), and  $I \rightarrow CV$  uses a stacked version of the analysis and assessment data (*i.e.*, all of the training data; approach 3) to impute missing values. Following CV, our analyses strictly implement approach 1 to engage with missing values in the testing data.

## 4 | SIMULATED EXPERIMENTS

The goal of the current simulation study was to assess empirical differences between  $CV \cup I$  and  $I \rightarrow CV$  following a published protocol.<sup>16</sup> Our primary objective was to measure and compare how well each strategy (1) approximated a model's generalization error and (2) selected parameters (both for imputation and modeling) that would maximize downstream model accuracy. To

complete item (1), we assessed estimation of the external  $R^2$ . We used bias, variance, and root-mean-squared error (RMSE) to quantify estimation accuracy. The RMSE provides an overall assessment of estimation accuracy that depends on both bias and variance. To complete item (2), we compared the performance (*i.e.*, the external  $R^2$ ) of downstream models whose tuning parameters were selected using  $CV \cup I$  versus  $I \rightarrow CV$ .

#### 4.1 | Data-generating mechanisms

Consider the linear regression model, where a continuous outcome vector  $\mathbf{y} = \{y_1, y_2, \dots, y_N\}$  is generated by a linear combination of predictor variables  $\mathbf{X} = [\mathbf{x}_1, \mathbf{x}_2, \dots, \mathbf{x}_p]$ . This functional relationship is often expressed as

$$\mathbf{Y} = \mathbf{X}\beta + \varepsilon,$$

where  $\beta$  is a  $p \times 1$  vector of regression coefficients and  $\varepsilon$  is a normally distributed  $N \times 1$  zero-mean random error vector. In practice,  $\mathbf{X}$  often has some ‘junk’ variables that are not related to the outcome. To make our simulations similar to applied settings, we generated normally distributed variables that had no relation to the simulated outcome. We fixed the number of true predictor variables at 10, the standard error of  $\varepsilon$  at 1, and set  $\beta = [-1.00, -0.78, -0.56, -0.33, -0.11, 0.11, 0.33, 0.56, 0.78, 1.00]$  throughout the simulation study. Columns of  $\mathbf{X}$  were generated from a multivariate normal distribution with a first order autoregressive correlation structure. Specifically, the correlation between columns  $\mathbf{x}_i$  and  $\mathbf{x}_j$  was  $\rho^{|i-j|}$ , where  $\rho$  was set to 3/4 throughout the study. We applied this design to generate a training set of varying size (100, 500, 1000, or 5000) along with an external validation set comprising 10,000 observations in each simulated replicate.

##### Data generation scenarios

We created three data-generation ‘scenarios’. In scenario 1, the observed data are independent and identically distributed (iid). In scenario 2, the data are iid conditional on an observed grouping variable. A total of 11 groups are formed, one in the validation set and the remaining 10 in the training set. Each group is characterized by a randomly generated mean value for its predictor variables. During CV, the observed groups are separated into ten folds to mimic the prediction of outcomes in a population with different characteristics. Scenario 3 is identical to scenario 2 except that the grouping variable is latent. Consequently, CV does not break the observed groups into separate folds for scenario 3.

##### Amputating data

We applied the `ampute` function from the `mi` R package to generate missing values in simulated data.<sup>17</sup> In each replicate, 90% of observations comprised at least one missing value. We designated up to  $p$  MD patterns randomly in each simulation replicate, where  $p$  is the number of non-outcome columns in the simulated data. A MD pattern indicates which of the  $p$  predictor variables are set to missing. For each MD pattern, the number of missing variables was randomly set to an integer ranging from 1 to  $p/2$ .

This procedure usually induced missing values in 30-50% of the data. When data were MAR, we applied the default method for the `ampute` function (`ampute.default.weights`) to induce missingness based on the observed variables. Throughout the experiment, we applied the same missing patterns and MD mechanism in the training set and the external validation set.

## Modeling procedure

We applied  $k$ -nearest-neighbor imputation to handle MD and least absolute shrinkage and selection operator (LASSO) regression to develop prediction functions throughout the simulated experiments.<sup>18</sup> The LASSO model is an appropriate model for these simulations since all data were generated with linear, independent effects. Nearest neighbor aggregation based on Gower's distance was used to form imputed values in the training and testing set and also in the analysis and assessment set for  $CV \cup I$ .<sup>19</sup> We created one imputed dataset for each  $k \in \{1, 2, \dots, 35\}$ . We selected a value for the regularization parameter  $\lambda$  in each imputed dataset, separately, using 10-fold CV (*i.e.*, `cv.glmnet`).<sup>20</sup> The  $\lambda$  value selected was the one that minimized the model's cross-validated RMSE. The value of  $k$  that minimized cross-validated RMSE was used to impute the entire training set prior to fitting a final `cv.glmnet` model.

## Analysis plan

We varied the scenario (1, 2, or 3; described above), missing mechanism (MCAR or MAR), ratio of predictor variables to junk variables (1:1, 1:4, and 1:49; junk variables have no relationship to the simulated outcome), and the number of training observations ( $N = 100, 500, 1,000, 5,000$ ). We present results for each of 72 settings determined by these parameters and also provide overall summary statistics for scenarios 1, 2, and 3 when data are MCAR and MAR (*i.e.*, aggregating over training sample size and predictor to noise ratio). In each simulation replication, we computed the true external  $R^2$  in the validation set for each potential value of nearest neighbors (*i.e.*,  $k \in \{1, 2, \dots, 35\}$ ). We also estimated external  $R^2$  for each value of  $k$  using  $CV \cup I$  and  $I \rightarrow CV$ , separately, to evaluate how well these CV procedures estimated the true external  $R^2$ . We assessed the difference between estimated external  $R^2$  according to  $CV \cup I$  and  $I \rightarrow CV$  as well as the bias, variance, and root-mean-squared error (RMSE) of these estimates. Last, we investigated the accuracy of downstream models when  $CV \cup I$  and  $I \rightarrow CV$  were applied to select the number of neighbors to use for imputation and then the regularization parameter for a penalized regression model.

## 4.2 | Results

Overall, a total of 143,253 out of 144,000 (99.5%) simulation replicates were completed over a span of 52,979 computing hours. Incomplete replicates were not analyzed, as these were replicates where at least one of the amputation, imputation, or prediction models did not converge. Across all replicates, the median (25<sup>th</sup>-75<sup>th</sup> percentile) number of seconds used to form imputed data using  $CV \cup I$  and  $I \rightarrow CV$  were 27 (11 - 225) and 2.3 (0.69 - 22), respectively, a ratio of 11 (9.8 - 14). Using the full imputed training set, the median (25<sup>th</sup>-75<sup>th</sup> percentile) number of seconds needed to tune `glmnet` models using CV and fit a final model

to the full training set was 4.8 (3.3 - 55), verifying our earlier claim that complex imputation procedures often require more time than modeling procedures.

Across all scenarios, the mean external  $R^2$  ranged from 0.233 to 0.443 (**Table 1**). External  $R^2$  values increased with larger training set size and higher ratio of predictor variables to junk variables. Notably, the mean external  $R^2$  values in scenario 1 were uniformly greater than corresponding mean external  $R^2$  values in scenarios 2 and 3, and the maximum difference between mean external  $R^2$  values in scenario 2 versus scenario 3 was 0.0003. The mean absolute difference between external  $R^2$  estimates using  $CV \cup I$  and  $I \rightarrow CV$  shrunk towards zero as the size of the training set increased (**Table 2**). The differences between  $CV \cup I$  and  $I \rightarrow CV$  were lowest in scenario 1 and greatest in scenario 2. These patterns were also present in visual depictions of external  $R^2$  portrayed as a function of  $k$  neighbors (**Figure 3**).

### Bias, variance, and RMSE

For scenario 1, the overall bias of  $R^2$  estimates under MCAR using  $CV \cup I$  was -0.00136 versus 0.00233 using  $I \rightarrow CV$  (**Table 3**). When the data were MAR, the overall biases were -0.00080 for  $CV \cup I$  versus 0.00300 for  $I \rightarrow CV$ . In scenarios 2 and 3, the bias of  $CV \cup I$  was lower than that of  $I \rightarrow CV$ , and  $I \rightarrow CV$  consistently provided overly optimistic error estimates. The overall standard deviation of  $R^2$  estimates was higher for  $CV \cup I$  versus  $I \rightarrow CV$  in all three scenarios and both missing data mechanisms. The difference in standard deviation was most pronounced in scenario 3 when data were MCAR (0.07314 [ $CV \cup I$ ] versus 0.06747 [ $I \rightarrow CV$ ]; **Table 4**). Despite the optimistic bias of  $I \rightarrow CV$  in scenario 2, the reduced variance of this approach lead to a lower overall RMSE for external  $R^2$  compared to  $CV \cup I$  (**Table 5**). When the data were MCAR in scenario 2,  $CV \cup I$  and  $I \rightarrow CV$  obtained RMSEs of 0.05738 and 0.05593, respectively. Similarly, when the data were MAR in scenario 2, overall RMSE values were 0.05751 and 0.05572.

### Downstream model performance

When  $CV \cup I$  and  $I \rightarrow CV$  were applied to select tuning parameters, the overall mean external  $R^2$  was higher using  $CV \cup I$  in 22 out of 78 comparisons (28%; **Table 6**). The greatest overall difference in mean  $R^2$  between downstream models occurred in scenario 1 when the data were MAR (absolute difference in model  $R^2$ : 0.00028; relative difference in model  $R^2$ : 0.07%).

## 5 | REAL DATA EXPERIMENTS

The goal of the current resampling study was to repeat comparisons summarized in Section 4 between  $CV \cup I$  and  $I \rightarrow CV$  using real, publicly accessible data. A secondary objective was to assess how much results would change if different modeling strategies were applied.



## Ames, Iowa housing data

The data we use in this resampling study describe the sale of individual residential property in Ames, Iowa from 2006 to 2010. The entire set contains 2930 observations and 80 variables (23 nominal, 23 ordinal, 14 discrete, and 20 continuous) that can be leveraged to predict the sale price of homes.<sup>21</sup> We used a cleaned version of the Ames data for our own analyses by applying the `make_ames()` function, available in the `AmesHousing` R package.<sup>22</sup> We also log-transformed the skewed sale price outcome.

## Analysis plan

We conducted a resampling study where the Ames housing data was randomly split into training ( $N = 2198$ , 75%) and testing ( $N = 732$ , 25%) sets in each of 5,000 iterations. In each resampling replicate, we implemented two separate modeling strategies to develop prediction functions using the training set: (1) un-penalized linear regression and (2) random forests.<sup>23</sup> We also implemented two imputation strategies: (1) nearest neighbor imputation using 1, 2, ..., 35 neighbors and (2) mean and mode imputation for numeric and nominal variables, respectively. In addition to imputation, data were pre-processed by lumping values in discrete variables into an ‘other’ category if the value accounted for less than 10% of the observed values. Both CV techniques (*i.e.*,  $CV \cup I$  and  $I \rightarrow CV$ ) were implemented to estimate the external generalization error of the linear regression and random forest models when nearest neighbor imputation was applied.

## Amputating data

The training and testing data were amputated in the same manner using four prototypes of missingness. The prototypes were characterized by having missing values for all variables describing the house (1) lot and garage, (2) longitude and latitude, (3) basement and year built, and (4) overall quality and general above ground square footage. We restricted our prediction models to consider only the 30 predictor variables belonging to at least one of these missing prototypes.

## 5.1 | Results

A total of 5000 out of 5000 (100%) resampling replicates were completed over a span of 927 computing hours. Across all replicates, the mean number of minutes used to form imputed data using  $CV \cup I$  and  $I \rightarrow CV$  were 10 and 1.0, respectively. The mean (standard deviation) external  $R^2$  for the linear regression and random forest models were 0.7641 (0.0265) and 0.8088 (0.0205), respectively. Overall, both CV techniques slightly over and under estimated the external  $R^2$  value when linear regression and random forests were applied, respectively. For linear regression, the mean (standard deviation) bias was -0.0015 (0.0343) and -0.0021 (0.0343) for  $CV \cup I$  and  $I \rightarrow CV$ , respectively. The standard deviations of error estimates were 0.0115 and 0.0115, respectively. For random forests, the mean (standard deviation) bias was 0.0009 (0.0255) and 0.0005 (0.0255) for  $CV \cup I$  and  $I \rightarrow CV$ , respectively. The standard deviation of error estimates were 0.0084 and 0.0086, respectively.

When  $CV \circ I$  and  $I \rightarrow CV$  were applied to select the number of neighbors used for imputation, downstream linear models obtained an external  $R^2$  of 0.7657 (0.0260) and 0.7659 (0.0259), respectively. Similarly, downstream random forests obtained an external  $R^2$  of 0.8097 (0.0202) and 0.8096 (0.0201), respectively. As a reference point, the mean (standard deviation) downstream external  $R^2$  when imputation to the mean was applied was 0.7329 (0.0283) and 0.8127 (0.0199) using linear regression and random forests, respectively. Overall, the computational resources required to implement  $CV \circ I$  were substantially higher than  $I \rightarrow CV$ , and the difference in downstream external  $R^2$  was  $<0.0001$  (Figure 4).

## Interpretation

The use of  $CV \circ I$  versus  $I \rightarrow CV$  resulted in a mean relative change of 0.0262% and -0.0220% in downstream model performance for linear models and random forests, respectively. These shifts in model performance come at the cost of a  $\nu$ -fold increase in the amount of computational resources allocated to handle MD. While improvements on the order of hundredths of a percentage may be notable for select analytic scenarios, these shifts in model performance may not be relevant in the majority of supervised learning analyses. In the latter case, it seems unsupervised imputation before CV can allow for pragmatic handling of MD without sacrificing the integrity of CV.

## 6 | DISCUSSION AND RECOMMENDATIONS

We demonstrated empirical properties of  $CV \circ I$  and  $I \rightarrow CV$  using nearest-neighbor imputation prior to applying regression and random forest models. We selected these methods because they have been studied thoroughly and are widely used in applied settings. In simulated experiments, we generated outcomes using linear effects without interaction. We also studied three broad scenarios that were relevant to CV: Scenario 1 was an ideal setting where  $I \rightarrow CV$  and  $CV \circ I$  should have provided almost identical estimates of generalization error. Scenarios 2 and 3 were meant to test whether  $I \rightarrow CV$  produced biased estimates of generalization error because in settings where  $I \rightarrow CV$  clearly did not mimic the final application of a trained model to an external validation set. Remarkably, despite its bias in scenario 2, the reduction in variance of  $R^2$  estimates using  $I \rightarrow CV$  lead to a lower overall RMSE compared to  $CV \circ I$ . Downstream model performance was consistently superior when  $CV \circ I$  was used instead of  $I \rightarrow CV$ . However, the increase in performance was smaller than 1% relative change (maximum overall relative difference in external  $R^2$ : 0.07%). While this difference is very small, it may be possible to find a different generative scenario where the difference is larger. Throughout our analysis,  $I \rightarrow CV$  required less computation time than  $CV \circ I$  by a factor of roughly  $\nu$ , the number of folds employed by CV.

Unsupervised imputation has two interesting characteristics relevant to predictive modeling. First, it allows for imputation of testing data without requiring observed outcome values in those data. This characteristic is ideal for deploying prediction models in the real world, where outcome values are almost always unknown at the time of prediction (otherwise, why would

we be predicting them?). Second, as the current analysis has shown, unsupervised imputation can be applied before CV in select scenarios without inducing overly optimistic estimates of model error. The benefits of this approach include (1) reduced computational overhead, (2) reduced variance in model error estimates, and (3) little difference in the performance of downstream models. If investigators are confident that training and testing data are identically distributed or are primarily concerned with selecting optimal tuning parameters,  $I \rightarrow CV$  may be an extremely valuable workflow to implement. However, the drawbacks of  $I \rightarrow CV$  include increased bias for model estimation, particularly in settings similar to scenario 2 (described in Section 4.1). If investigators are primarily interested in estimating model error without bias and cannot rule out the possibility that testing data are drawn from a different population or distribution than their training data, the current study suggests  $CV \circ I$  should be applied instead of  $I \rightarrow CV$ . However, it is worth noting that almost all prediction modeling decisions require balancing bias and variance to optimize precision. Our results do not indicate any strong difference between  $I \rightarrow CV$  and  $CV \circ I$  with regard to precision (*i.e.*, RMSE, see **Table 5**). We suspect that in most prediction modeling applications, precision rather than bias is of primary interest.

The current study has several strengths. We implemented computational experiments using simulated and real data. We included different data-generation mechanisms, different modeling procedures, different MD patterns, and different modeling strategies to ensure our results generalized to several common analytical settings. We examined a wide variety of metrics to assess the benefits and weaknesses of applying  $I \rightarrow CV$  versus  $CV \circ I$ . We used the *ipa* R package to conduct unsupervised imputation throughout our analyses and disseminate both the R package and all code used for the current analysis on the first author's GitHub. Each of these supplemental components ensure that our work is easily reproduced and disseminated. There are also some gaps in the current study that can be filled by future work. We investigated  $v$ -fold CV in the current analysis. Future research may assess whether these results generalize to other forms of data-splitting such as Monte-Carlo CV or bootstrap CV. Because MNAR data present challenges that may not be overcome by imputation alone, we did not include simulations for MNAR data. Whether the current study's findings generalize to settings with MNAR data remains an interesting, unanswered question. Last, the current study has applied  $k$ -nearest neighbor imputation throughout. As many other types of imputation procedures have been established, there are numerous extensions of the current analysis that may explore whether our results hold when other imputation approaches are implemented.

**TABLE 1** True external  $R^2$  mean (standard deviation) values for the modeling technique that is internally assessed using  $CV \cup I$  and  $I \rightarrow CV$ . Descriptions of scenarios 1, 2, and 3 are provided in Section 4.1. All table values are scaled by 100 for convenience

N	Scenario 1		Scenario 2		Scenario 3	
	MAR	MCAR	MAR	MCAR	MAR	MCAR
<b>10 predictors, 10 junk</b>						
100	37.8 (3.44)	37.7 (3.43)	33.5 (6.53)	33.3 (6.69)	33.5 (6.56)	33.3 (6.69)
500	42.8 (2.91)	42.7 (2.93)	40.0 (4.97)	39.8 (5.03)	40.0 (4.98)	39.8 (4.99)
1,000	43.5 (2.94)	43.4 (2.95)	40.9 (4.75)	40.7 (4.80)	40.9 (4.75)	40.7 (4.79)
5,000	44.3 (3.01)	44.2 (3.01)	42.0 (4.46)	41.8 (4.61)	42.0 (4.46)	41.8 (4.59)
<b>10 predictors, 40 junk</b>						
100	34.6 (3.83)	34.5 (3.75)	30.6 (6.39)	30.6 (6.23)	30.6 (6.44)	30.5 (6.28)
500	40.6 (2.77)	40.6 (2.77)	38.2 (4.67)	38.2 (4.73)	38.2 (4.68)	38.2 (4.61)
1,000	41.5 (2.73)	41.5 (2.74)	39.3 (4.68)	39.2 (4.67)	39.3 (4.67)	39.2 (4.68)
5,000	42.6 (2.75)	42.6 (2.75)	40.6 (4.11)	40.5 (4.17)	40.6 (4.11)	40.5 (4.20)
<b>10 predictors, 490 junk</b>						
100	27.5 (5.11)	27.6 (5.03)	23.3 (6.58)	23.3 (6.47)	23.3 (6.51)	23.3 (6.47)
500	37.6 (2.93)	37.6 (2.92)	35.8 (4.09)	35.9 (4.02)	35.8 (4.09)	35.9 (4.03)
1,000	38.8 (2.83)	38.7 (2.83)	37.2 (4.21)	37.2 (4.22)	37.2 (4.22)	37.2 (4.23)
5,000	39.8 (2.78)	39.8 (2.78)	38.5 (4.06)	38.5 (4.01)	38.5 (4.05)	38.5 (4.03)
<b>Overall</b>						
—	39.3 (5.51)	39.2 (5.46)	36.7 (7.18)	36.6 (7.16)	36.7 (7.17)	36.6 (7.14)

**TABLE 2** Mean (standard deviation) absolute differences in estimates of external  $R^2$  between  $CV \cup I$  and  $I \rightarrow CV$ . Descriptions of scenarios 1, 2, and 3 are provided in Section 4.1. All table values are scaled by 100 for convenience

N	Scenario 1		Scenario 2		Scenario 3	
	MAR	MCAR	MAR	MCAR	MAR	MCAR
<b>10 predictors, 10 junk</b>						
100	1.21 (1.02)	1.19 (0.99)	2.43 (1.76)	2.44 (1.79)	2.05 (1.64)	2.07 (1.67)
500	0.32 (0.29)	0.33 (0.29)	1.34 (1.13)	1.39 (1.19)	0.97 (1.05)	1.00 (1.09)
1,000	0.21 (0.20)	0.21 (0.20)	1.22 (1.15)	1.28 (1.21)	0.87 (1.07)	0.90 (1.12)
5,000	0.09 (0.08)	0.09 (0.09)	1.04 (0.92)	1.09 (1.01)	0.69 (0.89)	0.71 (0.95)
<b>10 predictors, 40 junk</b>						
100	1.49 (1.27)	1.47 (1.27)	2.71 (1.97)	2.67 (1.91)	2.38 (1.89)	2.35 (1.80)
500	0.33 (0.29)	0.34 (0.29)	1.40 (1.23)	1.45 (1.29)	1.01 (1.11)	1.04 (1.17)
1,000	0.22 (0.20)	0.22 (0.20)	1.27 (1.16)	1.33 (1.23)	0.90 (1.09)	0.93 (1.15)
5,000	0.09 (0.09)	0.10 (0.09)	1.15 (1.06)	1.23 (1.16)	0.77 (1.01)	0.80 (1.08)
<b>10 predictors, 490 junk</b>						
100	2.09 (1.73)	2.02 (1.67)	3.01 (2.18)	3.01 (2.14)	2.75 (2.04)	2.79 (2.05)
500	0.34 (0.29)	0.34 (0.29)	1.21 (1.15)	1.21 (1.14)	0.88 (1.05)	0.89 (1.05)
1,000	0.21 (0.19)	0.21 (0.19)	1.16 (1.16)	1.15 (1.15)	0.81 (1.04)	0.79 (1.02)
5,000	0.09 (0.08)	0.09 (0.08)	1.12 (1.11)	1.18 (1.17)	0.74 (1.03)	0.77 (1.08)
<b>Overall</b>						
—	0.56 (0.95)	0.55 (0.93)	1.59 (1.54)	1.62 (1.55)	1.23 (1.47)	1.25 (1.48)

**TABLE 3** Bias of external  $R^2$  estimates using  $CV \cup I$  and  $I \rightarrow CV$ . Descriptions of scenarios 1, 2, and 3 are provided in Section 4.1. All table values are scaled by 100 for convenience

N	Missing completely at random						Missing at random					
	Scenario 1		Scenario 2		Scenario 3		Scenario 1		Scenario 2		Scenario 3	
	$CV \cup I$	$I \rightarrow CV$	$CV \cup I$	$I \rightarrow CV$	$CV \cup I$	$I \rightarrow CV$	$CV \cup I$	$I \rightarrow CV$	$CV \cup I$	$I \rightarrow CV$	$CV \cup I$	$I \rightarrow CV$
<b>10 predictors, 10 junk</b>												
100	0.31	-0.48	1.06	-1.14	0.57	-1.21	0.51	-0.31	1.32	-0.85	0.76	-0.99
500	0.10	-0.02	0.26	-1.06	-0.20	-1.07	0.11	0.01	0.41	-0.86	-0.05	-0.90
1,000	0.09	0.04	0.20	-1.04	-0.26	-1.08	0.10	0.04	0.36	-0.83	-0.08	-0.87
5,000	0.00	0.00	0.11	-0.97	-0.32	-0.99	0.01	0.00	0.28	-0.75	-0.11	-0.77
<b>10 predictors, 40 junk</b>												
100	0.62	-0.52	1.02	-1.44	0.56	-1.56	0.80	-0.39	1.33	-1.17	0.82	-1.31
500	0.14	-0.01	0.36	-1.01	-0.09	-1.02	0.20	0.07	0.45	-0.88	-0.02	-0.91
1,000	0.04	-0.04	0.19	-1.11	-0.29	-1.15	0.06	-0.02	0.22	-1.02	-0.22	-1.04
5,000	0.06	0.03	0.16	-1.06	-0.31	-1.09	0.05	0.03	0.22	-0.92	-0.20	-0.94
<b>10 predictors, 490 junk</b>												
100	1.08	-0.74	1.44	-1.43	1.26	-1.39	1.02	-0.88	1.42	-1.46	1.15	-1.46
500	0.23	0.08	0.38	-0.68	0.02	-0.70	0.37	0.21	0.48	-0.58	0.11	-0.61
1,000	0.12	0.06	0.20	-0.87	-0.20	-0.88	0.31	0.24	0.38	-0.68	-0.01	-0.71
5,000	0.00	-0.02	0.26	-0.90	-0.20	-0.93	0.05	0.03	0.26	-0.85	-0.16	-0.86
<b>Overall</b>												
—	0.23	-0.14	0.47	-1.06	0.04	-1.09	0.30	-0.08	0.59	-0.90	0.16	-0.95

**TABLE 4** Standard deviation of external  $R^2$  estimates using  $CV \cup I$  and  $I \rightarrow CV$ . Descriptions of scenarios 1, 2, and 3 are provided in Section 4.1. All table values are scaled by 100 for convenience

N	Missing completely at random						Missing at random					
	Scenario 1		Scenario 2		Scenario 3		Scenario 1		Scenario 2		Scenario 3	
	$CV \cup I$	$I \rightarrow CV$	$CV \cup I$	$I \rightarrow CV$	$CV \cup I$	$I \rightarrow CV$	$CV \cup I$	$I \rightarrow CV$	$CV \cup I$	$I \rightarrow CV$	$CV \cup I$	$I \rightarrow CV$
<b>10 predictors, 10 junk</b>												
100	6.41	6.38	7.56	7.09	7.46	7.04	6.37	6.33	7.52	7.13	7.44	7.06
500	3.59	3.64	3.92	3.82	3.92	3.81	3.60	3.65	3.92	3.82	3.91	3.81
1,000	3.32	3.36	3.67	3.55	3.65	3.54	3.31	3.35	3.67	3.55	3.65	3.54
5,000	3.08	3.10	3.39	3.29	3.40	3.28	3.07	3.09	3.37	3.28	3.38	3.27
<b>10 predictors, 40 junk</b>												
100	6.73	6.62	7.28	6.91	7.19	6.88	6.62	6.51	7.30	6.89	7.23	6.85
500	3.55	3.60	3.75	3.65	3.70	3.65	3.57	3.61	3.73	3.64	3.71	3.65
1,000	3.14	3.17	3.31	3.22	3.32	3.21	3.16	3.19	3.28	3.21	3.31	3.21
5,000	2.82	2.84	3.02	2.89	3.02	2.89	2.82	2.84	2.99	2.88	3.00	2.88
<b>10 predictors, 490 junk</b>												
100	7.66	7.43	7.89	7.51	7.84	7.50	7.52	7.31	8.00	7.61	7.96	7.56
500	3.69	3.73	3.85	3.78	3.86	3.78	3.72	3.75	3.85	3.77	3.87	3.78
1,000	3.26	3.29	3.30	3.21	3.30	3.21	3.27	3.29	3.33	3.24	3.34	3.25
5,000	2.86	2.87	2.95	2.87	2.98	2.87	2.87	2.88	2.94	2.87	2.97	2.88
<b>Overall</b>												
—	6.50	6.12	7.31	6.77	7.31	6.75	6.52	6.12	7.34	6.79	7.33	6.76

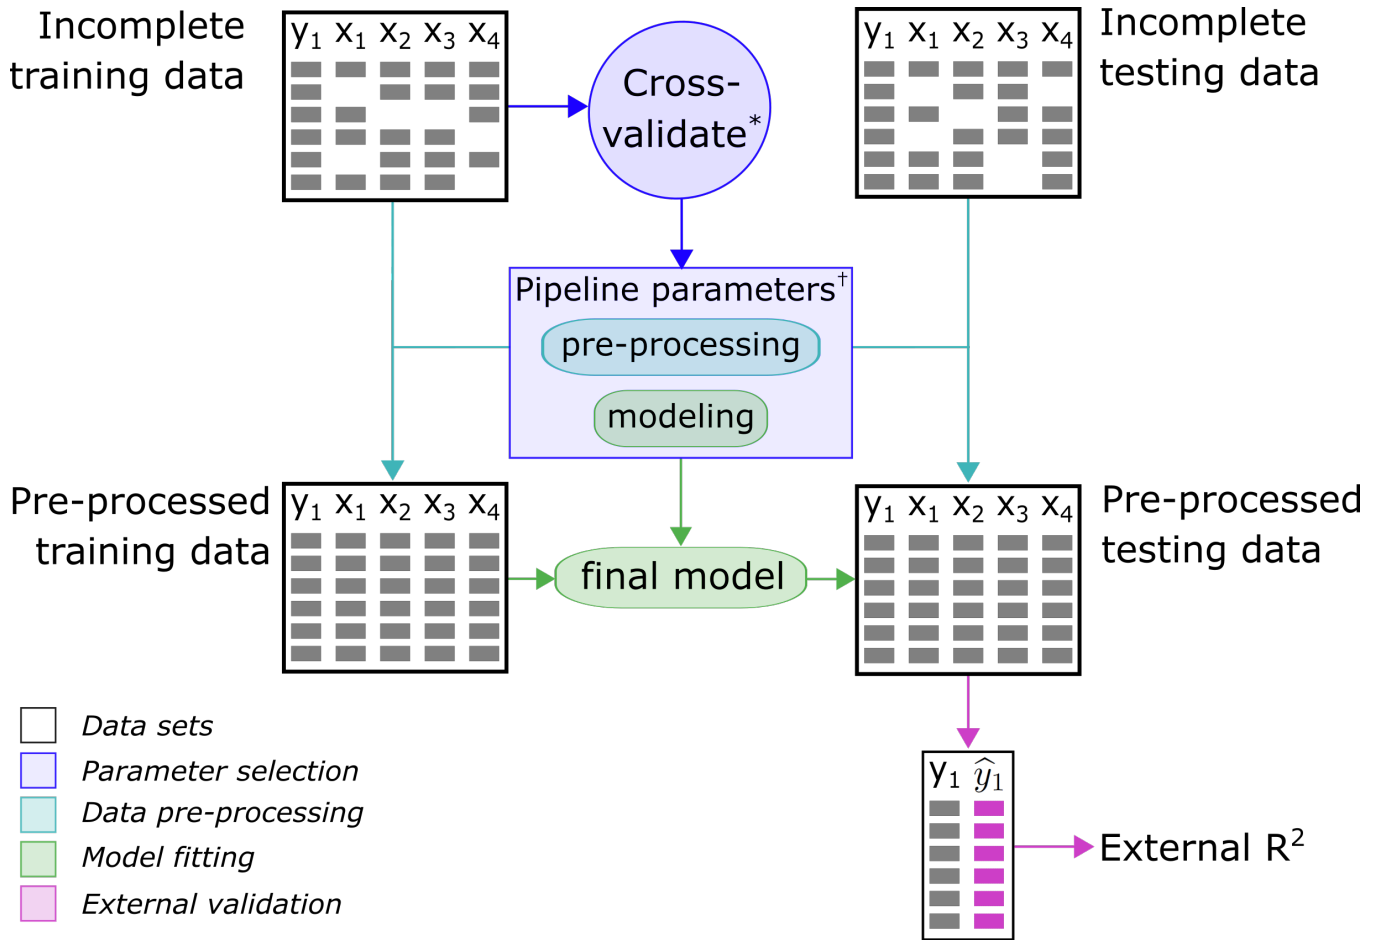
**TABLE 5** Root-mean-squared error of external  $R^2$  estimates using  $CV \cup I$  and  $I \rightarrow CV$ . Descriptions of scenarios 1, 2, and 3 are provided in Section 4.1. All table values are scaled by 100 for convenience

N	Missing completely at random						Missing at random					
	Scenario 1		Scenario 2		Scenario 3		Scenario 1		Scenario 2		Scenario 3	
	$CV \cup I$	$I \rightarrow CV$	$CV \cup I$	$I \rightarrow CV$	$CV \cup I$	$I \rightarrow CV$	$CV \cup I$	$I \rightarrow CV$	$CV \cup I$	$I \rightarrow CV$	$CV \cup I$	$I \rightarrow CV$
<b>10 predictors, 10 junk</b>												
100	5.93	5.86	9.08	8.74	8.94	8.69	5.87	5.77	8.98	8.62	8.84	8.59
500	2.22	2.22	4.98	4.91	4.91	4.86	2.22	2.22	4.94	4.84	4.92	4.85
1,000	1.64	1.64	4.31	4.20	4.25	4.22	1.64	1.64	4.25	4.11	4.19	4.12
5,000	0.82	0.82	3.67	3.63	3.64	3.61	0.82	0.82	3.48	3.40	3.46	3.40
<b>10 predictors, 40 junk</b>												
100	6.36	6.26	8.77	8.53	8.72	8.59	6.36	6.23	8.97	8.64	8.92	8.68
500	2.36	2.36	4.79	4.69	4.61	4.59	2.34	2.33	4.74	4.61	4.66	4.62
1,000	1.66	1.66	4.39	4.37	4.39	4.39	1.69	1.69	4.42	4.37	4.39	4.37
5,000	0.82	0.82	3.58	3.48	3.59	3.52	0.81	0.81	3.45	3.35	3.44	3.36
<b>10 predictors, 490 junk</b>												
100	7.38	7.31	8.93	8.79	8.90	8.80	7.35	7.35	9.06	8.89	8.91	8.80
500	2.43	2.42	4.17	4.01	4.14	4.02	2.45	2.43	4.24	4.05	4.21	4.07
1,000	1.76	1.75	3.91	3.78	3.89	3.80	1.76	1.75	3.92	3.72	3.87	3.75
5,000	0.85	0.85	3.38	3.24	3.37	3.27	0.86	0.86	3.42	3.28	3.38	3.28
<b>Overall</b>												
—	3.62	3.58	5.74	5.59	5.68	5.59	3.61	3.57	5.75	5.57	5.68	5.56

**TABLE 6** Mean external  $R^2$  when  $CV \cup I$  and  $I \rightarrow CV$  were applied to tune the number of neighbors used for imputation. Descriptions of scenarios 1, 2, and 3 are provided in Section 4.1. All table values are scaled by 100 for convenience

N	Missing completely at random						Missing at random					
	Scenario 1		Scenario 2		Scenario 3		Scenario 1		Scenario 2		Scenario 3	
	$CV \cup I$	$I \rightarrow CV$	$CV \cup I$	$I \rightarrow CV$	$CV \cup I$	$I \rightarrow CV$	$CV \cup I$	$I \rightarrow CV$	$CV \cup I$	$I \rightarrow CV$	$CV \cup I$	$I \rightarrow CV$
<b>10 predictors, 10 junk</b>												
100	38.1	38.1	33.9	34.0	34.0	34.0	38.3	38.3	34.2	34.2	34.1	34.1
500	43.8	43.8	40.7	40.8	40.8	40.8	43.9	43.9	41.0	41.0	41.0	41.0
1,000	44.7	44.8	41.9	41.9	41.9	41.9	44.8	44.8	42.1	42.1	42.1	42.1
5,000	45.7	45.7	43.3	43.3	43.3	43.3	45.8	45.8	43.5	43.5	43.5	43.5
<b>10 predictors, 40 junk</b>												
100	35.1	35.0	31.1	31.2	31.1	31.1	35.2	35.0	31.1	31.0	31.1	31.1
500	41.8	41.8	39.1	39.2	39.2	39.2	41.8	41.8	39.2	39.2	39.2	39.2
1,000	42.8	42.8	40.3	40.3	40.3	40.3	42.8	42.8	40.4	40.4	40.4	40.4
5,000	44.0	44.0	41.9	41.9	41.9	41.9	44.0	44.0	42.0	42.0	42.0	42.0
<b>10 predictors, 490 junk</b>												
100	28.8	28.5	24.0	23.7	24.0	23.7	28.6	28.4	24.1	23.8	24.2	23.9
500	39.1	39.1	37.2	37.3	37.2	37.3	39.1	39.1	37.2	37.3	37.2	37.2
1,000	40.4	40.4	38.7	38.7	38.7	38.7	40.4	40.4	38.7	38.7	38.7	38.7
5,000	41.5	41.5	40.1	40.1	40.1	40.1	41.5	41.5	40.1	40.1	40.1	40.1
<b>Overall</b>												
—	40.5	40.5	37.7	37.7	37.7	37.7	40.5	40.5	37.8	37.8	37.8	37.8

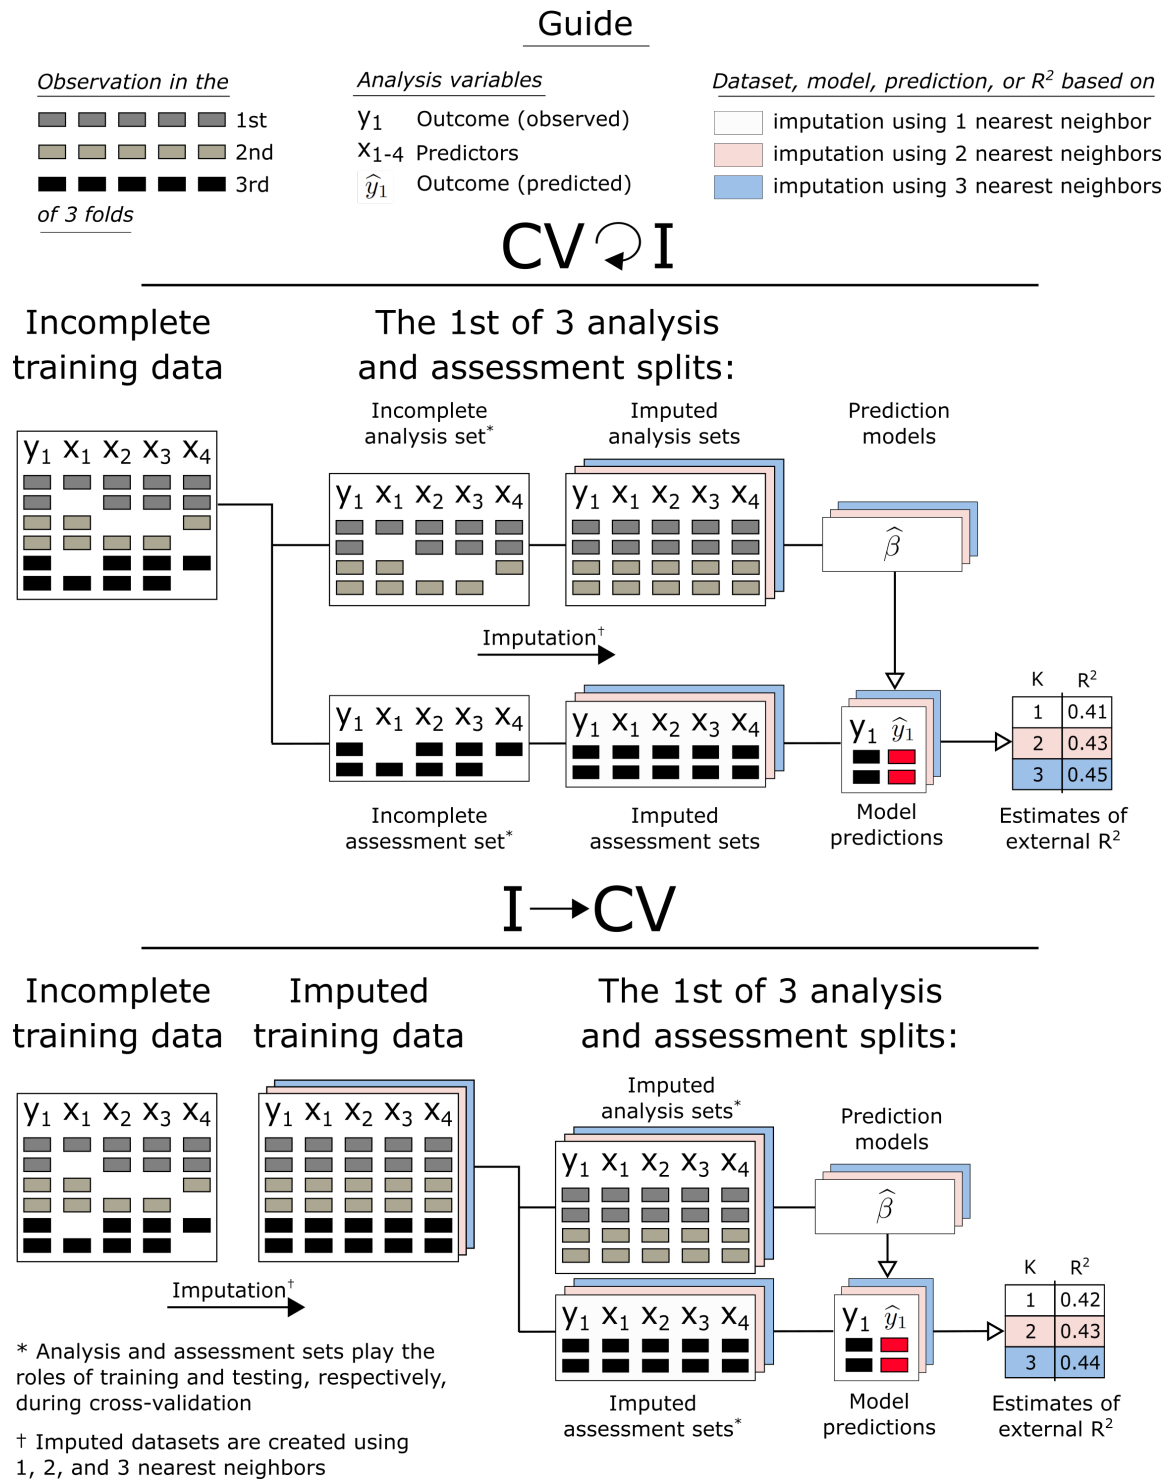





\* details on this analysis stage are provided in Figure 2

† pipeline parameters are used for both pre-processing and model fitting


**FIGURE 1** A workflow to develop and validate a pipeline modeling algorithm. Pipeline parameter values may be set apriori or determined empirically using cross validation. Once parameter values are fixed, a final model is developed by training the modeling pipeline using the training data. The final model is externally validated by assessing the accuracy of its predictions in the testing data.



**FIGURE 2** Workflows for cross validation (CV) incorporating imputation of missing values. The difference in the workflows is where imputation is performed. The standard workflow,  $CV \rightarrow I$ , imputes missing values during each replicate of CV. The experimental workflow,  $I \rightarrow CV$ , imputes missing values prior to CV. Critically,  $I \rightarrow CV$  means imputation happens once, whereas in  $CV \rightarrow I$  the imputation procedure occurs for each fold, adding computational time.



**FIGURE 3** External generalization error and internal estimates of generalization error using  $I \rightarrow CV$  and  $CV \cap I$ . The  $R^2$  values are plotted as a function of the number of nearest neighbors used to impute missing data, and the panel rows show results with 100, 500, 1000, and 5000 observations in the training data. The scenarios are described in Section 4.1. The peaks of all three curves consistently appear to be at the same number of neighbors. While  $I \rightarrow CV$  error estimates have a slight positive bias, as noted in section 4.2, they also have less variability than error estimates using  $CV \cap I$ .



**FIGURE 4** External generalization error (y-axis) and time required to impute missing values (x-axis) for the Ames housing data.  $I \rightarrow CV$  and  $CV \cap I$  were applied, separately, to select  $k$ , the number of nearest neighbors used to impute missing values, prior to fitting and externally validating a prediction model. While the median generalization error is practically equivalent regardless of which CV method was used, the time required for imputation is approximately 10 (*i.e.*, the number of folds in CV) times higher using  $CV \cap I$  versus  $I \rightarrow CV$ .

## References

1. Kuhn M, Johnson K, others . *Applied predictive modeling*. 26. Springer . 2013.
2. Arlot S, Celisse A, others . A survey of cross-validation procedures for model selection. *Statistics surveys* 2010; 4: 40–79.
3. Lang M, Binder M, Richter J, et al. mlr3: A modern object-oriented machine learning framework in R. *Journal of Open Source Software* 2019. doi: 10.21105/joss.01903
4. Hastie T, Tibshirani R, Friedman J. *The elements of statistical learning: data mining, inference, and prediction*. 2. Springer Science & Business Media . 2009.
5. Stekhoven DJ, Bühlmann P. MissForest—non-parametric missing value imputation for mixed-type data. *Bioinformatics* 2011; 28(1): 112–118.
6. Rubin DB. Inference and missing data. *Biometrika* 1976; 63(3): 581–592.
7. Twala B, Jones M, Hand DJ. Good methods for coping with missing data in decision trees. *Pattern Recognition Letters* 2008; 29(7): 950–956.
8. Twala B. An empirical comparison of techniques for handling incomplete data using decision trees. *Applied Artificial Intelligence* 2009; 23(5): 373–405.
9. Tang F, Ishwaran H. Random forest missing data algorithms. *Statistical Analysis and Data Mining: The ASA Data Science Journal* 2017; 10(6): 363–377.
10. Van Buuren S. *Flexible imputation of missing data*. CRC press . 2018.
11. Kuhn M, Johnson K. *Feature engineering and selection: A practical approach for predictive models*. CRC Press . 2019.
12. Jerez JM, Molina I, García-Laencina PJ, et al. Missing data imputation using statistical and machine learning methods in a real breast cancer problem. *Artificial intelligence in medicine* 2010; 50(2): 105–115.
13. Kuhn M, J S. *Tidy modeling with R*. Springer . 2020.
14. Neunhoeffler M, Sternberg S. How cross-validation can go wrong and what to do about it. *Political Analysis* 2019; 27(1): 101–106.
15. Hastie T, Mazumder R. softImpute: Matrix Completion via Iterative Soft-Thresholded SVD. 2015. R package version 1.4.
16. Morris TP, White IR, Crowther MJ. Using simulation studies to evaluate statistical methods. *Statistics in medicine* 2019; 38(11): 2074–2102.

17. van Buuren S, Groothuis-Oudshoorn K. mice: Multivariate Imputation by Chained Equations in R. *Journal of Statistical Software* 2011; 45(3): 1-67.
18. Tibshirani R. Regression shrinkage and selection via the lasso. *Journal of the Royal Statistical Society: Series B (Methodological)* 1996; 58(1): 267–288.
19. Gower JC. A general coefficient of similarity and some of its properties. *Biometrics* 1971: 857–871.
20. Friedman J, Hastie T, Tibshirani R. Regularization Paths for Generalized Linear Models via Coordinate Descent. *Journal of Statistical Software* 2010; 33(1): 1–22.
21. De Cock D. Ames, Iowa: Alternative to the Boston housing data as an end of semester regression project. *Journal of Statistics Education* 2011; 19(3).
22. Kuhn M. AmesHousing: The Ames Iowa Housing Data. 2017. R package version 0.0.3.
23. Breiman L. Random forests. *Machine learning* 2001; 45(1): 5–32.

